

THE BLUEPRINT

The PromptGrade Blueprint

A field guide to writing prompts that score 90+ on the
Prompt Efficiency Score™ — and ship better AI output every day.

8 PES Dimensions

12 Power Patterns

Daily Checklist

v1.0 · Premium Edition · 18 pages of actionable frameworks

Table of Contents

01	Why most prompts fail (and what 90+ prompts do differently)	p. 3
02	The Prompt Efficiency Score™ — the 8 dimensions explained	p. 4
03	The PES Blueprint — an 8-step framework you can run today	p. 6
04	12 Power Patterns — copy-paste templates with examples	p. 9
05	Model-specific tuning — GPT, Claude, Gemini, Llama	p. 13
06	Before / After — six real prompt rewrites	p. 15
07	Your daily 60-second prompt checklist	p. 17
08	Where to go from here	p. 18

Why Most Prompts Fail

Most people prompt LLMs the same way they Google: a few keywords, fingers crossed. The result is generic output, hallucinated facts, and hours of re-prompting. After analyzing tens of thousands of prompts through the PromptGrade scanner, a clear pattern emerged.

The 4 silent killers of prompt quality

- 1. Vague role** — no persona, expertise, or perspective specified.
- 2. Missing constraints** — no length, format, tone, or audience.
- 3. Zero context** — the model has to guess your situation.
- 4. No success criteria** — you can't measure 'good' if you didn't define it.

High-performing prompts treat the model like a senior contractor: brief them like a pro, give them the inputs they need, and define what 'done' looks like. The next chapters show you exactly how.

The Prompt Efficiency Score™

Every prompt analyzed by PromptGrade is scored across 8 weighted dimensions. Here's what each one measures and the lever that moves it the most.

Dimension	Measures	Highest-leverage fix
Clarity	Is the request unambiguous?	Replace pronouns with nouns. Define jargon. State one objective.
Specificity	Are the details quantified?	Add numbers: length, count, dates, percentages, examples.
Context	Does the model know your situation?	Provide the who/why/where. One sentence of background is worth a thousand retries.
Role	Is a persona or expertise assigned?	'Act as a senior X with 10 years experience in Y.'
Format	Is the output structure defined?	Specify markdown, JSON schema, table columns, or section headers.
Constraints	Are the boundaries explicit?	State must / must-not, length caps, tone, reading level, banned phrases.
Examples	Did you show, not tell?	Add 1–3 few-shot examples of desired output for instant calibration.
Success Criteria	How will you grade the answer?	List 2–3 acceptance checks the model can self-verify against.

Scoring math, simplified

Each dimension is scored 0–10 and weighted by impact: Context, Constraints, and Specificity carry the heaviest weights because they have the biggest measured effect on output quality. A prompt that scores 90+ has explicit signal in **all eight** dimensions — there are no free passes.

The PES Blueprint — 8 Steps

Run this sequence top-to-bottom and you will produce a 90+ prompt in under 4 minutes. Keep this page open while you write.

Step 01 · Set the Role

Open with one line: 'You are a [seniority] [role] with deep expertise in [domain].' Force the model into the right knowledge slice.

Step 02 · State the Goal

One sentence: 'I want to [verb] [object] for [audience] so that [outcome].' If you can't write it, you're not ready to prompt.

Step 03 · Drop the Context

Paste the situation, the data, or the constraints the model has no way to know. Two paragraphs max.

Step 04 · Define Inputs

List the variables the model should treat as inputs (URLs, names, numbers). Wrap them in <> so they're unambiguous.

Step 05 · Specify Format

Markdown? JSON? A table with columns A, B, C? Be ruthless. Provide a skeleton if structure matters.

Step 06 · Set Constraints

Length, tone, reading level, banned words, required citations. State both must and must-not.

Step 07 · Add 1–2 Examples

Show one ideal answer and (optionally) one bad one labelled 'Avoid'. Few-shot beats fine-tuning for one-offs.

Step 08 · Demand Self-Check

Close with: 'Before answering, verify your draft against these criteria: [list]. Revise once, then output.'

Pro move

Save the 8-step skeleton as a snippet in your IDE or notes app. 90% of your prompts can be assembled by filling in the blanks in under 3 minutes.

12 Power Patterns

Twelve battle-tested templates you can copy verbatim. Swap the <> variables and ship.

PATTERN 01 · Role + Goal + Constraint

Use when: You need any decent baseline answer.

Template

Act as a <<role>>. Your goal is to <<goal>>. Constraints: <<constraints>>.

Example

Act as a senior B2B copywriter. Your goal is to draft 3 cold-email subject lines for a CFO audience. Constraints: ≤ 7 words each, no emojis, no clickbait.

PATTERN 02 · Chain-of-Thought Trigger

Use when: Multi-step reasoning, math, or analysis.

Template

Think step-by-step. First, list assumptions. Second, work through the logic. Third, give the final answer in <answer> tags.

Example

Used to solve a complex pricing decision with explicit assumptions surfaced.

PATTERN 03 · Few-Shot Calibration

Use when: You want a very specific tone or shape.

Template

Here are 2 examples of the style I want.\nExample 1: ...\nExample 2: ...\nNow produce one for: <<input>>.

Example

Used to mimic a brand voice across 50 product descriptions.

12 Power Patterns

PATTERN 04 · JSON Contract

Use when: You need parseable output for code.

Template

Reply with ONLY valid JSON matching this schema: { "title": string, "tags": string[], "score": number 0-100 }. No prose.

Example

Used to extract structured metadata from articles for a pipeline.

PATTERN 05 · Critic + Reviser

Use when: Quality matters more than speed.

Template

Draft an answer. Then critique it as a harsh senior <<role>> would. Then rewrite incorporating the critique. Output only the final version.

Example

Used to polish executive memos before publishing.

PATTERN 06 · Anti-Hallucination Guard

Use when: Factual accuracy is non-negotiable.

Template

If you are not at least 90% confident in a fact, write [UNVERIFIED] before it. If you don't know, say 'I don't know' – don't guess.

Example

Used in legal research drafts to flag claims for human review.

12 Power Patterns

PATTERN 07 · Audience Translator

Use when: Same content, different reader.

Template

Rewrite the following for <<audience>> at a <<reading level>> reading level. Preserve all facts.

Example

Used to convert a technical changelog into a customer-facing release note.

PATTERN 08 · Compare and Decide

Use when: Choosing between options.

Template

Compare <<A>> vs <> on: <<criteria>>. Output a table, then a recommendation with a one-paragraph rationale.

Example

Used to pick between two SaaS vendors with explicit trade-offs.

PATTERN 09 · Outline-Then-Expand

Use when: Long-form writing.

Template

First produce a numbered outline. Wait for me to approve. Then expand each section into <<word count>> words.

Example

Used for 2000-word thought-leadership posts.

12 Power Patterns

PATTERN 10 · Persona Interview

Use when: Discovery and research.

Template

Roleplay as <<persona>>. Answer the next 5 questions in-character with realistic objections and emotion.

Example

Used to pressure-test product positioning with simulated buyers.

PATTERN 11 · Constraint Stacking

Use when: When the first answer is too generic.

Template

Same task, now add these constraints: no buzzwords, max 120 words, must include one concrete example, must end with a question.

Example

Used to sharpen LinkedIn drafts after a weak first pass.

PATTERN 12 · Self-Eval Rubric

Use when: You want the model to grade itself.

Template

After answering, score your response on Clarity, Specificity, and Usefulness (0-10 each) and explain any score below 8 in one line.

Example

Used to flag weak answers for human review without re-prompting.

Model-Specific Tuning

Different frontier models reward different prompt shapes. Calibrate, don't copy-paste blindly.

GPT-4o / GPT-5 — Loves structure & explicit JSON contracts.

Use system messages aggressively. Wrap variables in <<brackets>>. Ask for JSON when chaining.

Claude (Sonnet/Opus) — Best with XML tags & long context.

Use <instructions>, <context>, <output> tags. Claude follows nested structure faithfully and excels at multi-doc reasoning.

Gemini (1.5/2.0) — Strong on multimodal & retrieval.

Front-load context (Gemini weights early tokens). Use explicit step numbering. Provide image/PDF inputs when available.

Llama 3 / Open models — Reward few-shot examples heavily.

Always include 1–3 examples. Be more directive — open models drift more without strong scaffolding.

Universal truth

No matter the model, the highest-leverage edits are still: assign a role, define success criteria, and show one example. Master those before you chase model-specific tricks.

Before / After Rewrites

Vague brief → Tight brief

Before	After
Write me a marketing email about our new feature.	Act as a senior B2B email copywriter. Draft a 120-word product announcement email for engineering managers at Series B SaaS companies. The new feature is auto-scaling for our queue product (saves 30% infra cost). Tone: confident, no hype, no emojis. Output: subject line + body in markdown. End with a one-line CTA to the changelog.

No format → JSON contract

Before	After
Pull the action items from this meeting transcript.	Extract action items from the transcript. Reply with ONLY valid JSON: [{"owner": string, "task": string, "due": ISO-date null}]. If owner is unclear, set to "unassigned". No prose.

No success criteria → Self-checked

Before	After
Summarize this article.	Summarize the article in ≤ 120 words for a busy executive. Before finalizing, verify: (a) every claim is in the source, (b) no jargon without a 5-word gloss, (c) ends with the 1 implication for our business. Revise once if any check fails.

The 60-Second Checklist

Print this. Tape it to your monitor. Run it on every prompt before you press send.

- **Role assigned** — Did I tell the model who to be?
- **Goal stated** — Did I write one sentence describing the outcome?
- **Context given** — Does the model know my situation and audience?
- **Inputs labelled** — Are my variables clearly marked?
- **Format defined** — Did I specify the structure of the output?
- **Constraints set** — Length, tone, must/must-not — all covered?
- **Example provided** — Did I show at least one ideal answer?
- **Self-check requested** — Does the prompt end with a verification step?

If your prompt fails 3+ checks...

Don't send it. Spend 60 more seconds fixing the cheapest fix first (usually Format or Role).
Your future self will thank you when you don't have to re-prompt 5 times.

Where to Go From Here

You now have the framework, the patterns, and the checklist. The fastest way to internalize all of it is to **run every prompt you write through the PromptGrade scanner** for the next 14 days. You'll see your score climb from the 50s into the 90s — and your AI output will follow.

Free, unlimited scans on the web app

Head to promptgrade.app and paste any prompt to get an instant PES score, line-by-line annotations, and a rewritten version optimized for your target model. No install. No credit card.

Want everything unlocked — multi-model rewrites, prompt chains, masterclass, and the workspace analytics dashboard? Upgrade to PromptGrade Premium for a one-time \$7.99. Lifetime access.

© 2026 PromptGrade™ · promptgrade.app · Made for prompt engineers who ship.